Noisy Channel Language Model Prompting for Few-shot Text Classification

Sewon Min, Mike Lewis, Hannaneh Hajishirzi, Luke Zettlemoyer
University of Washington, Meta AI, Allen Institute for AI
sewon@cs.washington.edu / @sewon__min

PAPER

Background

LM Prompting (Brown et al 2020):
using a frozen LM for a downstream task

- No or very limited parameter updates
- High variance, low worst-case accuracy

Method

- ✓ Zero-shot inference
- ✓ In-context learning
- ✓ Ensemble-based In-context learning
- ✓ Prompt tuning (Lester et al 2021)

[Baselines]

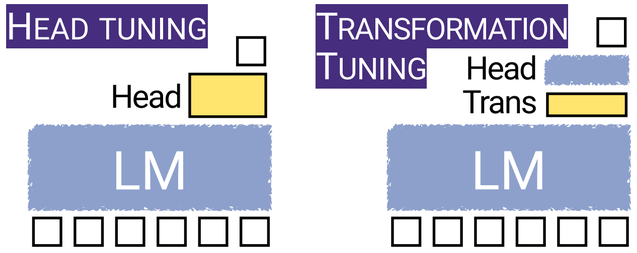

DIRECT: $P(y \mid x)$

x Why are boolean values capitalized in Python? → LM → It is about Computer & Internet. y

CHANNEL: $P(x \mid y)P(y) \propto P(x \mid y)$

y It is about Computer & Internet. → LM → Why are boolean values capitalized in Python? x

$$P(y \mid x) = \frac{P(x \mid y)P(y)}{P(x)} \propto P(x \mid y)P(y)$$

Why does it work?

- ✓ Better at few-shot (Ng and Jordan, 2002, Ding and Gimpel, 2019)
- ✓ More robust to distribution shift (Yogtama et al. 2017, Lewis and Fan, 2018)
- ✓ Required to predict the entire input (Lewis & Fan, 2018)
- ✓ Insensitive to label distributions

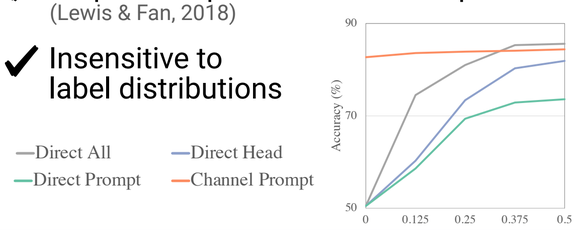

Experiments

16-shot data
4 x labels
5 x data
4 x train

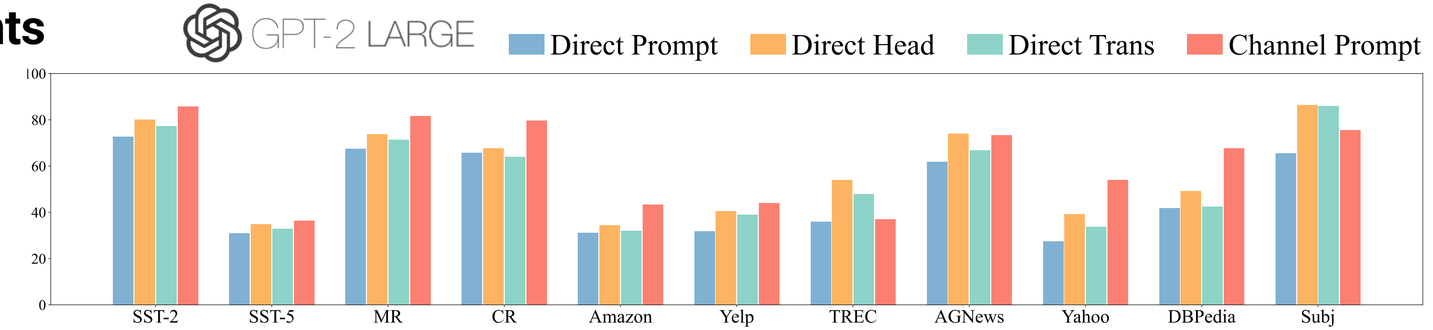

- ✓ Head tuning is a powerful baseline
- ✓ Direct models suffer from high variance & low worst-case accuracy
- ✓ Channel models have significantly lower variance and higher worst-case accuracy → better performance on average

Ablations in the paper:
Channel is better with …

- *Small k*
- *Large # of labels*
- *Imbalanced data*
- *Distribution shift*

CODE & DATA